

Lighter: Configuration-Driven Deep Learning

Ibrahim Hadzic^{1,2¶}, Suraj Pai^{2,3,4}, Keno Bressemer^{5,6}, Borek Foldyna¹, and Hugo JWL Aerts^{2,3,4}

¹ Cardiovascular Imaging Research Center, Massachusetts General Hospital, Harvard Medical School, United States of America ² Radiology and Nuclear Medicine, CARIM & GROW, Maastricht University, The Netherlands ³ Artificial Intelligence in Medicine (AIM) Program, Mass General Brigham, Harvard Medical School, Harvard Institutes of Medicine, United States of America ⁴ Department of Radiation Oncology, Brigham and Women's Hospital, Dana-Farber Cancer Institute, Harvard Medical School, United States of America ⁵ Technical University of Munich, School of Medicine and Health, Klinikum rechts der Isar, TUM University Hospital, Germany ⁶ Department of Cardiovascular Radiology and Nuclear Medicine, Technical University of Munich, School of Medicine and Health, German Heart Center, TUM University Hospital, Germany ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Lighter is an open-source Python [framework](#) for deep learning research that builds upon PyTorch Lightning ([Falcon & The PyTorch Lightning team, 2019](#)) and the [MONAI Bundle configuration](#) ([Cardoso et al., 2022](#)). With its declarative YAML-based configuration system that is both transparent and self-documenting, Lighter aims to streamline deep learning research. Researchers define experimental protocols—including neural network architectures, optimization strategies, data pipelines, and evaluation metrics—through structured configuration files, which effectively decouple scientific hypotheses from implementation details. This separation reduces boilerplate code while preserving complete experimental control. Lighter enables reproducibility through comprehensive configuration snapshots that document all experimental parameters and dependencies. A modular adapter system and support for project-specific extensions ensure the extensibility of the framework enabling researchers to implement specialized methodologies without modifying core framework components. By abstracting the engineering complexities of deep learning experimentation, Lighter allows researchers to focus on scientific innovation, accelerate hypothesis testing, and facilitate rigorous validation of research findings across application domains.

Statement of Need

Lighter is designed to address several key challenges in deep learning experimentation:

- Boilerplate Code:** Writing code for training loops, data loading, metric calculations, and experiment setups is repetitive and can vary greatly between projects. *Lighter abstracts these repetitive tasks, exposing only the components that differ across projects.*
- Experiment Management:** Handling numerous hyperparameters and configurations across various experiments can become cumbersome and error-prone. *Lighter offers organized configuration through YAML files, providing a centralized record of all experiment parameters.*
- Reproducibility:** Reproducing experiments from different implementations can be challenging. *Lighter's self-contained configuration files serve as comprehensive documentation, facilitating the exact recreation of experimental setups.*
- Collaboration:** Collaborating on experiments often requires understanding complex

codebases. *Lighter enhances collaboration by using standardized configurations, making it easier to share and reuse experiment setups within and across research teams.*

5. **Slowed Iteration:** The cumulative effect of these challenges slows down the research iteration cycle. *Lighter accelerates iteration by streamlining the experiment setup process, allowing researchers to focus on core experiment choices without being bogged down by infrastructure concerns.*

Design

Lighter's architecture is built upon three fundamental components that work together to streamline deep learning experimentation:

1. **Config:** This component serves as the experiment's blueprint, parsing and validating YAML configuration files that comprehensively define all aspects of the experimental setup. Within these configuration files, researchers specify the System and Trainer parameters, creating a self-documenting record of the experiment.
2. **System:** The System orchestrates the main building blocks of an experiment: model, optimizer, scheduler, loss function, evaluation metrics, and dataloaders. It implements the logic controlling how these components interact during training, validation, testing, and inference phases, that is modifiable through [adapters](#).
3. **Trainer:** PyTorch Lightning's robust Trainer class handles the technical aspects of the training process. It manages advanced features such as distributed training across multiple GPUs, mixed precision computation for memory efficiency, and checkpoint management for experiment continuity. Lighter employs the Trainer to execute the experimental protocol defined by the System.

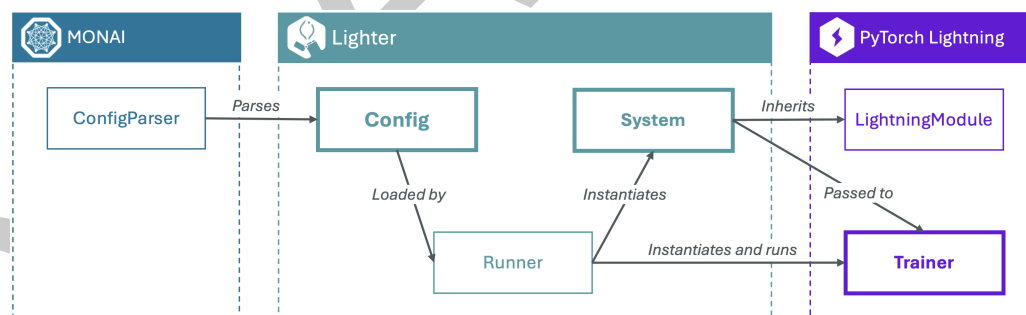


Figure 1: Lighter Overview. Lighter revolves around three main components – Trainer, System and Config, which contains the definition for the former two. Config leverages MONAI's ConfigParser for parsing the user-defined YAML configuration files, and its features are used by Runner to instantiate the System and Trainer. Trainer is used directly from PyTorch Lightning, whereas System inherits from LightningModule, ensuring its compatibility with Trainer while implementing a logic generalizable to any task or type of data. Finally, Runner runs the paired Trainer and System for a particular stage (e.g., fit or test).

System

The System component in Lighter serves as a comprehensive abstraction layer that encapsulates the essential experimental elements—including neural network architectures, optimization strategies, learning rate schedulers, loss functions, evaluation metrics, and data pipelines. This component implements a generalized data flow paradigm that orchestrates interactions between these elements during the experimental lifecycle. A distinguishing feature of the System is its configurable nature through the [adapter](#) mechanism, which provides the flexibility required to address diverse research tasks without architectural modifications. The systematic flow of

information through the System is illustrated in Figure 2, demonstrating how data traverses from input through model inference to evaluation and logging.

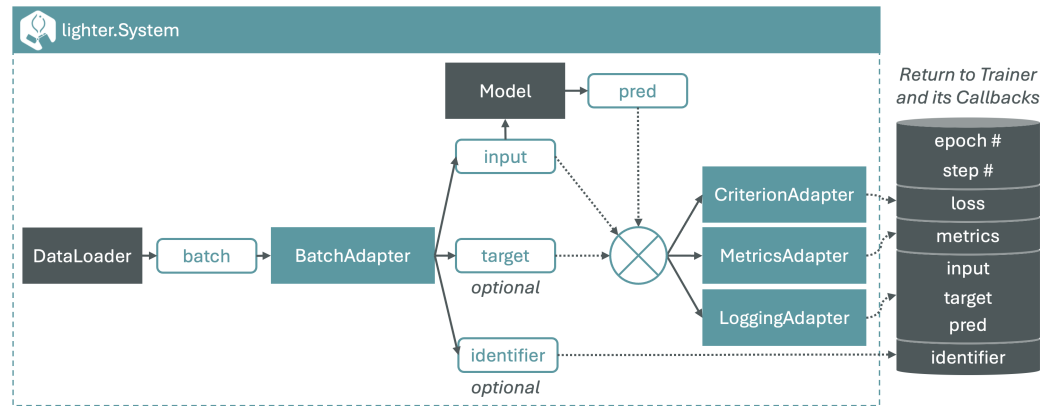


Figure 2: Flowchart of the `Lighter.System`. A batch from the `DataLoader` is processed by `BatchAdapter` to extract input, target (optional), and identifier (optional). The `Model` generates `pred` (predictions) from the input. `CriterionAdapter` and `MetricsAdapter` compute loss and metrics, respectively, by applying optional transformations and adapting arguments for the loss and metric functions. Argument adaptation reorders or names inputs; for example, if a loss function expects `loss_fn(predictions, ground_truth)`, the `CriterionAdapter` maps `pred` to `predictions` and `target` to `ground_truth`. `LoggingAdapter` prepares data for logging. Results, including loss, metrics, and processed data, are returned to the `Trainer`.

Adaptability Through Modular Design

Lighter achieves task-agnostic flexibility through two key concepts: adapters and project-specific module integration.

Adapters

The adapter pattern implements a transformation layer between core system components, enabling customized data flow between the dataloader, criterion, metrics computation, and logging subsystems (Figure 2). This abstraction allows researchers to modify component interactions without altering the underlying framework. Consider the following configuration excerpt that demonstrates the adapter's capability to transform prediction outputs and remap function arguments:

```

adapters:
  train:
    criterion:
      _target_: lighter.adapters.CriterionAdapter
      pred_transforms: # Apply sigmoid activation to predictions
      _target_: torch.sigmoid
      pred_argument: 0 # Pass 'pred' to criterion's 1st arg
      target_argument: 1 # Pass 'target' to criterion's 2nd arg

```

As a result, Lighter can execute a wide range of deep learning tasks, from classification and segmentation to self-supervised learning, without requiring changes to the core framework.

Project-specific modules

Lighter provides the integration of project-specific implementations through a modular project structure. Researchers can use their custom components—including novel architectures, specialized datasets, task-specific metrics, and domain-adapted transforms—within a structured

90 project directory. This organization promotes code reusability and maintains a clear separation
91 between framework functionality and project-specific implementations.

92 For example, given a project folder `joss_project` with the following structure:

```
93 joss_project
94 |__ __init__.py
95 |__ models/
96 |    |__ __init__.py
97 |    |__ mlp.py
```

98 This folder will be imported as a module named `project`, which can then be used to reference
99 the components defined within it:

```
project: /path/to/joss_project
system:
  model:
    _target_: project.models.mlp.MLP
    input_size: 784
    num_classes: 10
```

100 Research Contributions That Use Lighter

101 Lighter has enabled advancements in medical imaging research:

- 102 ▪ Foundation model for cancer imaging biomarkers (Pai et al., 2024)
- 103 ▪ Vision Foundation Models for Computed Tomography (Pai et al., 2025)

104 Acknowledgments

105 We thank John Zielke for the idea to use adapter design pattern. We thank the MONAI team
106 (Wenqi Li, Nic Ma, Yun Liu, Eric Kerfoot) for their continuous support with features and
107 improvements related to MONAI Bundle.

108 References

- 109 Cardoso, M. J., Li, W., Brown, R., Ma, N., Kerfoot, E., Wang, Y., Murray, B., Myronenko, A.,
110 Zhao, C., Yang, D., Nath, V., He, Y., Xu, Z., Hatamizadeh, A., Zhu, W., Liu, Y., Zheng,
111 M., Tang, Y., Yang, I., ... Feng, A. (2022). *MONAI: An open-source framework for deep*
112 *learning in healthcare*. <https://doi.org/https://doi.org/10.48550/arXiv.2211.02701>
- 113 Falcon, W., & The PyTorch Lightning team. (2019). *PyTorch Lightning* (Version 1.4).
114 <https://doi.org/10.5281/zenodo.3828935>
- 115 Pai, S., Bontempi, D., Hadzic, I., Prudente, V., Sokač, M., Chaunzwa, T. L., Bernatz, S.,
116 Hosny, A., Mak, R. H., Birkbak, N. J., & Aerts, H. J. W. L. (2024). Foundation model for
117 cancer imaging biomarkers. *Nat. Mach. Intell.*, 6(3), 354–367.
- 118 Pai, S., Hadzic, I., Bontempi, D., Bressem, K., Kann, B. H., Fedorov, A., Mak, R. H., & Aerts,
119 H. J. W. L. (2025). *Vision foundation models for computed tomography*.